

Part 2: Weights



@NAMlab




www.szymanskilab.com



CEPLAS
Cluster of Excellence on Plant Sciences



Artificial Neural Networks




Vectorizing Forward Computation

For now let's ignore bias

$$W^{(1)} = \begin{matrix} \xrightarrow{\text{Inputs in layer } l-1} \\ \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} & w_{13}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} & w_{23}^{(1)} \end{bmatrix} \downarrow \text{Neurons in layer } l \end{matrix}$$

$w_{ij}^{(l)}$: weight from j^{th} neuron on layer l to the i^{th} neuron on layer $l+1$



$$X = \begin{matrix} \xrightarrow{\text{Features/Inputs}} \\ \begin{bmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \\ x_{41} & x_{42} & x_{43} \end{bmatrix} \downarrow \text{Observations} \end{matrix}$$

x_{ij} : i^{th} sample for feature j


Vectorizing Forward Computation

For now let's ignore bias


$$W^{(1)} \cdot X^T = \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} & w_{13}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} & w_{23}^{(1)} \end{bmatrix} \cdot \begin{bmatrix} x_{11} & x_{21} & x_{31} & x_{41} \\ x_{12} & x_{22} & x_{32} & x_{42} \\ x_{13} & x_{23} & x_{33} & x_{43} \end{bmatrix}$$

Observation Features/Inputs

$$= \begin{bmatrix} w_{11}^{(1)}x_{11} + w_{12}^{(1)}x_{12} + w_{13}^{(1)}x_{13} & \dots & w_{11}^{(1)}x_{41} + w_{12}^{(1)}x_{42} + w_{13}^{(1)}x_{43} \\ w_{21}^{(1)}x_{11} + w_{22}^{(1)}x_{12} + w_{23}^{(1)}x_{13} & \dots & w_{21}^{(1)}x_{41} + w_{22}^{(1)}x_{42} + w_{23}^{(1)}x_{43} \end{bmatrix}$$



Forward Computation

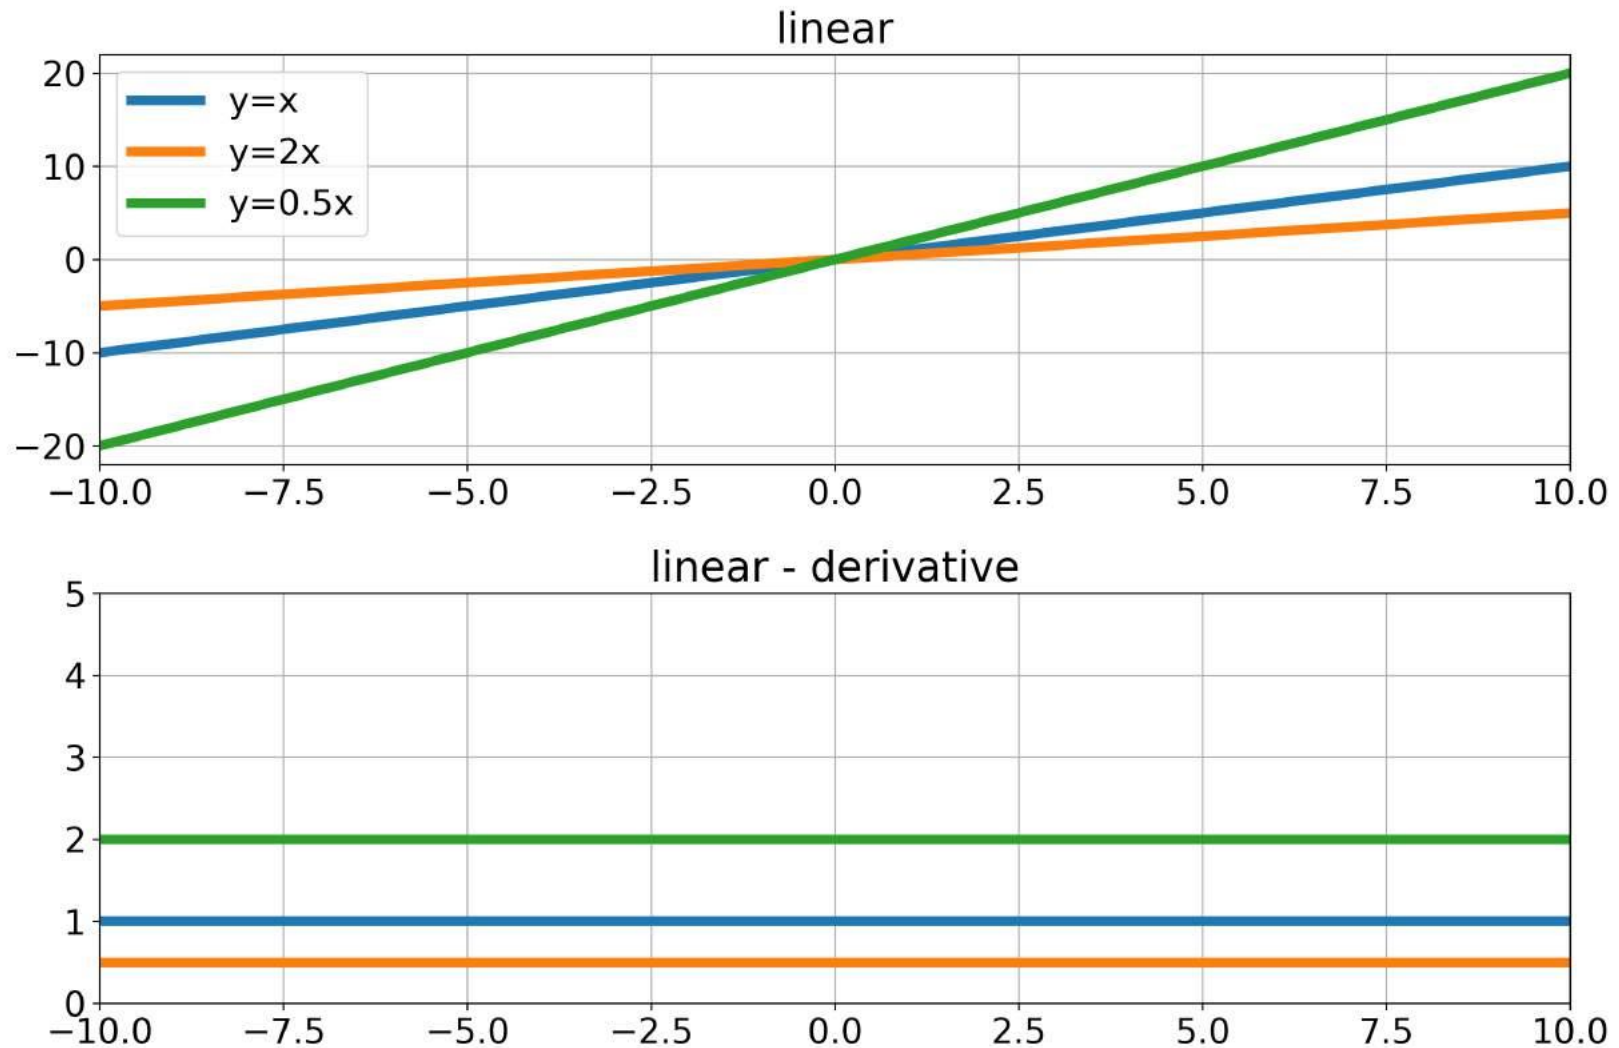


On every layer we multiply the output of the previous layer with the weights and apply activation function $\varphi(\cdot)$:

$$z^{(l)} = W^{(l)} X^T$$
$$a^{(l)} = \varphi(z^{(l)})$$

Note the $W^{(l)}$ contains also the bias

Linear Activation




Vanishing/Exploding Gradients

Reminder: Backpropagation uses the chain rule to compute gradients.

Let's examine gradient of the first layer:

$$\frac{\partial L}{\partial a} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial d} \frac{\partial d}{\partial c} \frac{\partial c}{\partial b} \frac{\partial b}{\partial a}$$


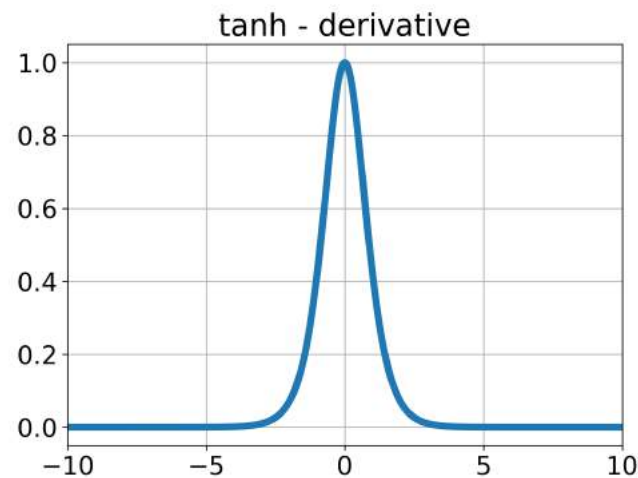






Activation Functions

$$\varphi(x) = \frac{1}{1 + e^{-x}}$$

$$\varphi(x) = \frac{e^{2x} - 1}{e^{2x} + 1}$$

$$\varphi(x) = \max(0, x)$$



Vanishing/Exploding Activations

To simplify assume linear activation function $\varphi(x) = x$ and 0 bias. Then the NN output will be:

$$\hat{y} = W^{(n)} W^{(n-1)} \dots \overbrace{W^{(2)} W^{(1)}}^{\text{second layer}} X^T$$

$\underbrace{\hspace{10em}}_{\text{first layer}}$

If we initialize weights randomly:


- $W \gg 1$

What will happen to the output of the NN?

Vectorizing Forward Computation

$$W^{(1)} = \begin{bmatrix} \text{Bias} & \text{Inputs in layer } l-1 \\ w_{10}^{(1)} & w_{11}^{(1)} & w_{12}^{(1)} & w_{13}^{(1)} \\ w_{20}^{(1)} & w_{21}^{(1)} & w_{22}^{(1)} & w_{23}^{(1)} \end{bmatrix} \begin{matrix} \downarrow \\ \text{Neurons in layer } l \end{matrix}$$

$w_{ij}^{(l)}$: weight from j^{th} neuron on layer l to the i^{th} neuron on layer $l+1$




$$X = \begin{bmatrix} \text{For Bias} & \text{Features/Inputs} \\ 1 & x_{11} & x_{12} & x_{13} \\ 1 & x_{21} & x_{22} & x_{23} \\ 1 & x_{31} & x_{32} & x_{33} \\ 1 & x_{41} & x_{42} & x_{43} \end{bmatrix} \begin{matrix} \downarrow \\ \text{Observations} \end{matrix}$$

x_{ij} : i^{th} sample for feature j

Weight Initialisation

Q: Why not initialize all the weights to 0?


$W=0$



Weight Initialisation (LeCun)

Initialise with small random number $N(0, \alpha)$
 $\alpha = 1e-2$

Works for small networks but leads to skewed activations for deeper networks

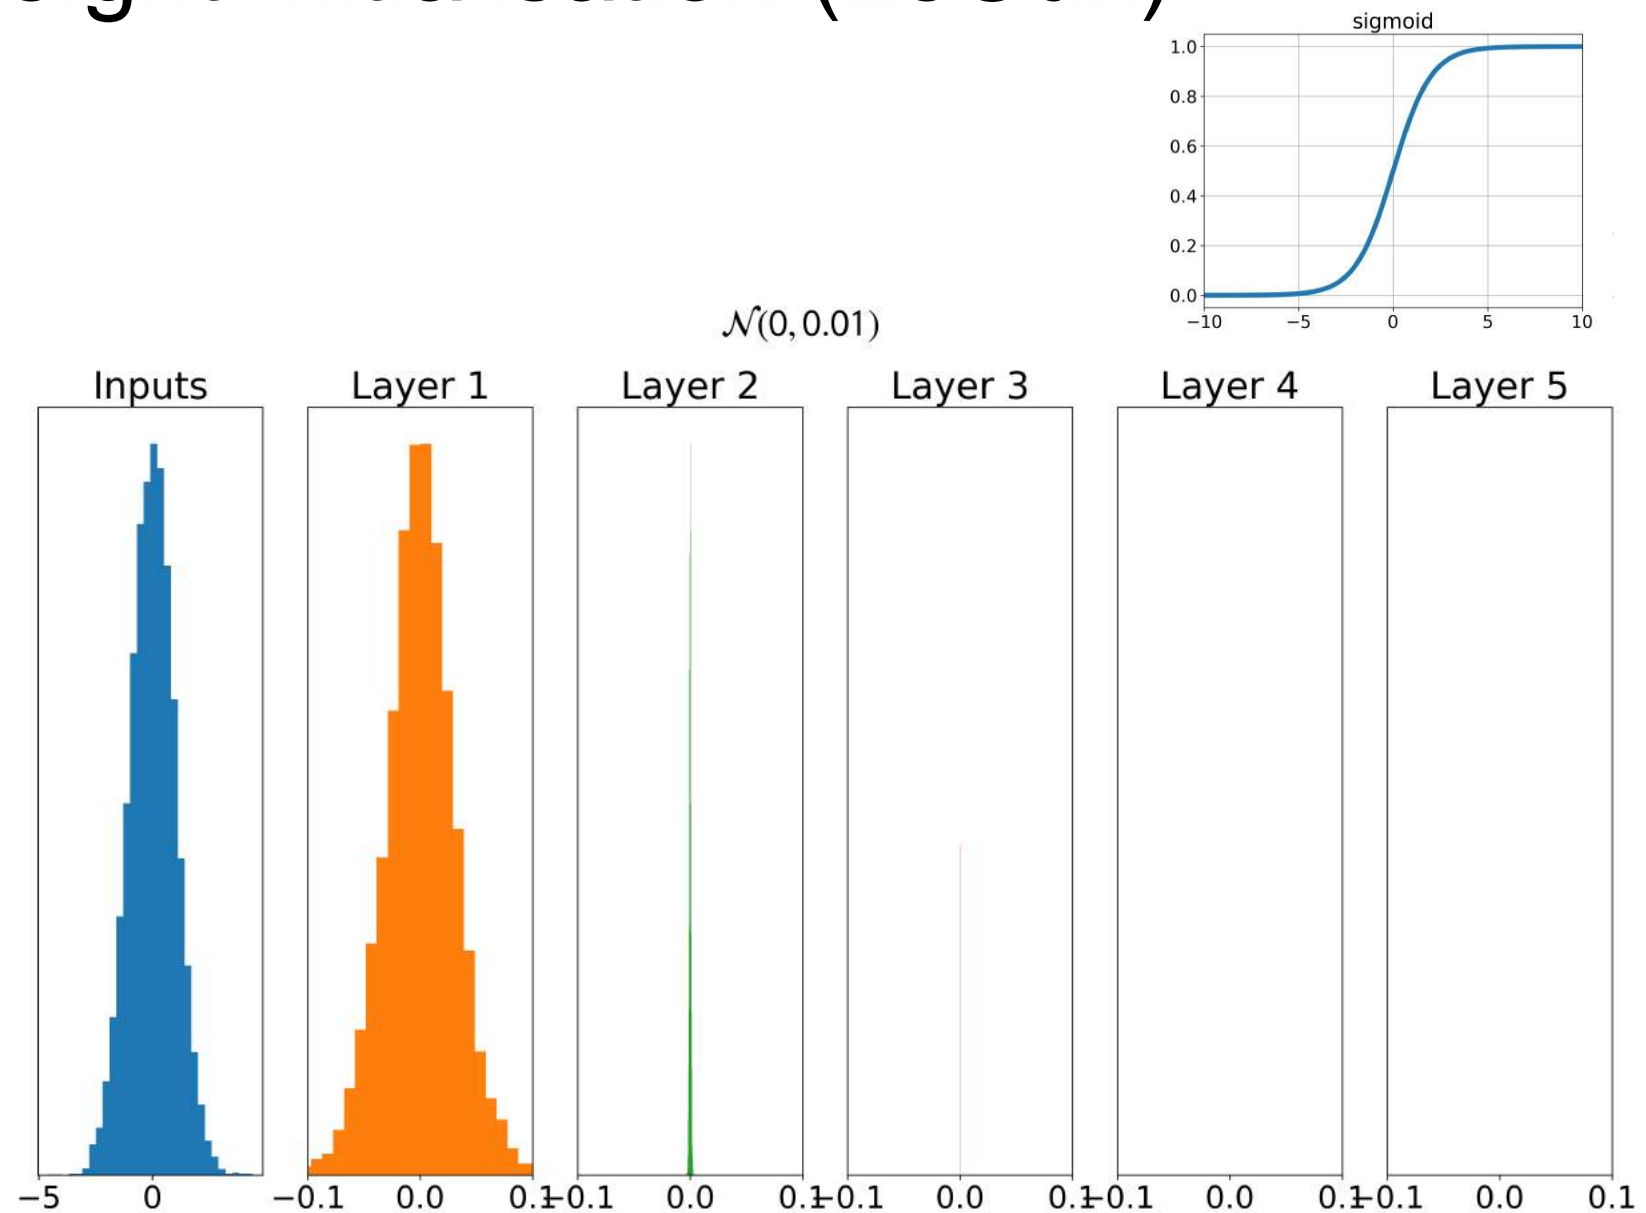


Weight Initialisation (LeCun)

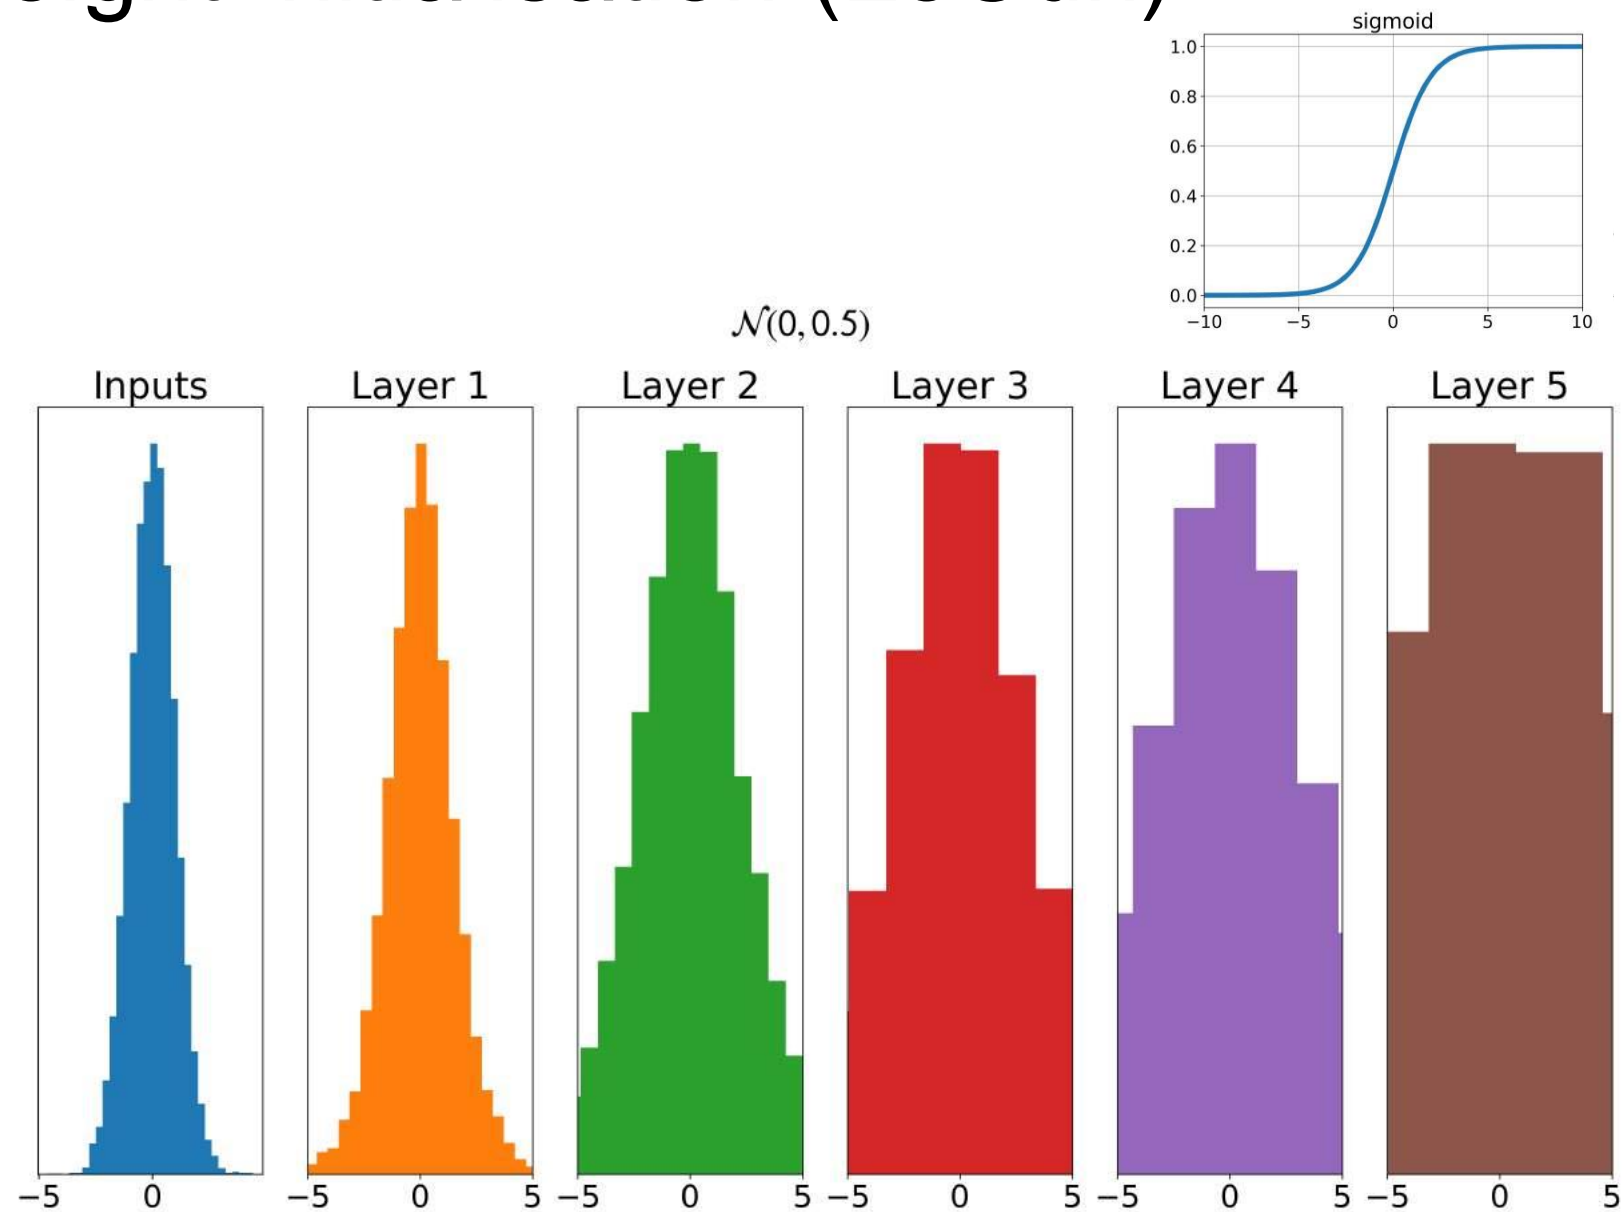
- Let's assume that we have 10-D input data x data normally distributed in $N(0, 1)$
- Let's create 5 layers with 10 neurons on each layer
- Let's initialise weights from using a normal distribution $N(0, 1e-2)$

What is the expected distribution of activations after a forward pass?

Weight Initialisation (LeCun)





Weight Initialisation (LeCun)



Xavier Initialisation

$$\sqrt{\frac{1}{\text{size}^{(l-1)}}}$$



$\mathcal{N}(0, 0.31622776601683794)$



Xavier with ReLU

$$\sqrt{\frac{1}{\text{size}^{(l-1)}}}$$

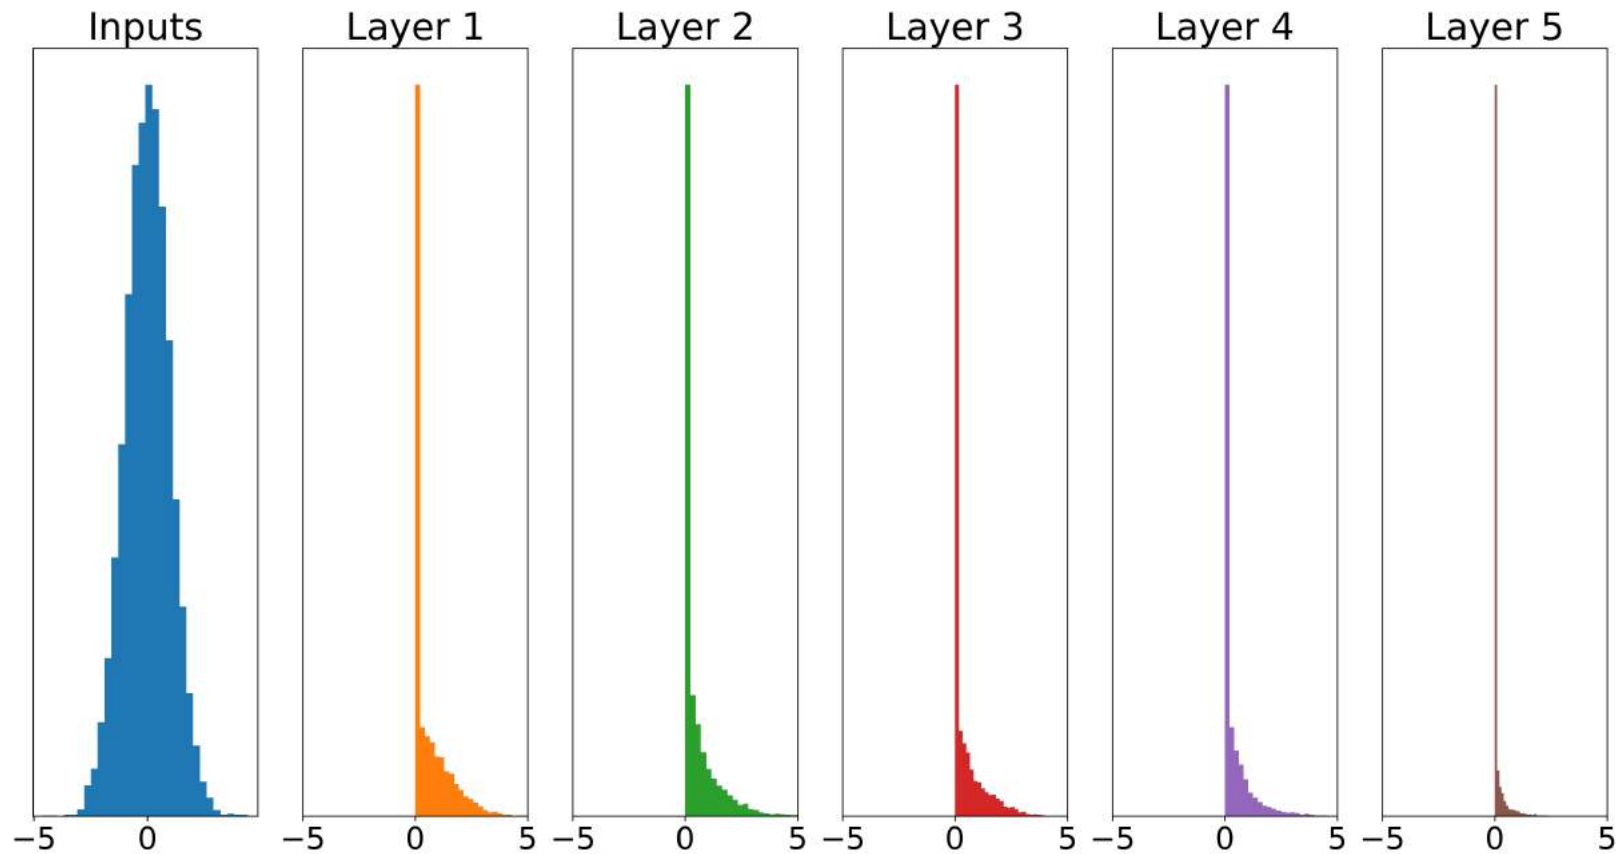

$\mathcal{N}(0, 0.31622776601683794)$



Kaiming Initialisation

$$\sqrt{\frac{2}{\text{size}^{(l-1)} + \text{size}^{(l)}}}$$

$\mathcal{N}(0, 0.4472135954999579)$



Batch Normalisation

If we need zero mean, unit variance in every layer why not normalize at every layer for every mini-batch?


$$\mathcal{B} = \{x_1 \dots x_m\}$$

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \varepsilon}}$$

Where $\mu_{\mathcal{B}}$ and $\sigma_{\mathcal{B}}^2$ the mean and variance of the minibatch

$$\mu_{\mathcal{B}} = \frac{1}{m} \sum_{i=1}^m x_i$$
$$\sigma_{\mathcal{B}}^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$$

Batch Normalisation



This way we constrain the layers and non-linear activations to operate in zero mean, unit variance input and this may not be desirable.

$$y_i = \gamma \hat{x}_i + \beta$$
$$BN_{\gamma, \beta}(x_i)$$

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \varepsilon}} \quad \gamma = \sqrt{\sigma_B^2 + \varepsilon} \quad \beta = \mu_B$$